


Budapest DB Meetup – 2018/Nov/13

Mapping Graph Queries to PostgreSQL

Gábor Szárnyas, József Marton,
Márton Elekes, János Benjamin Antal



Magyar Tudományos
Akadémia



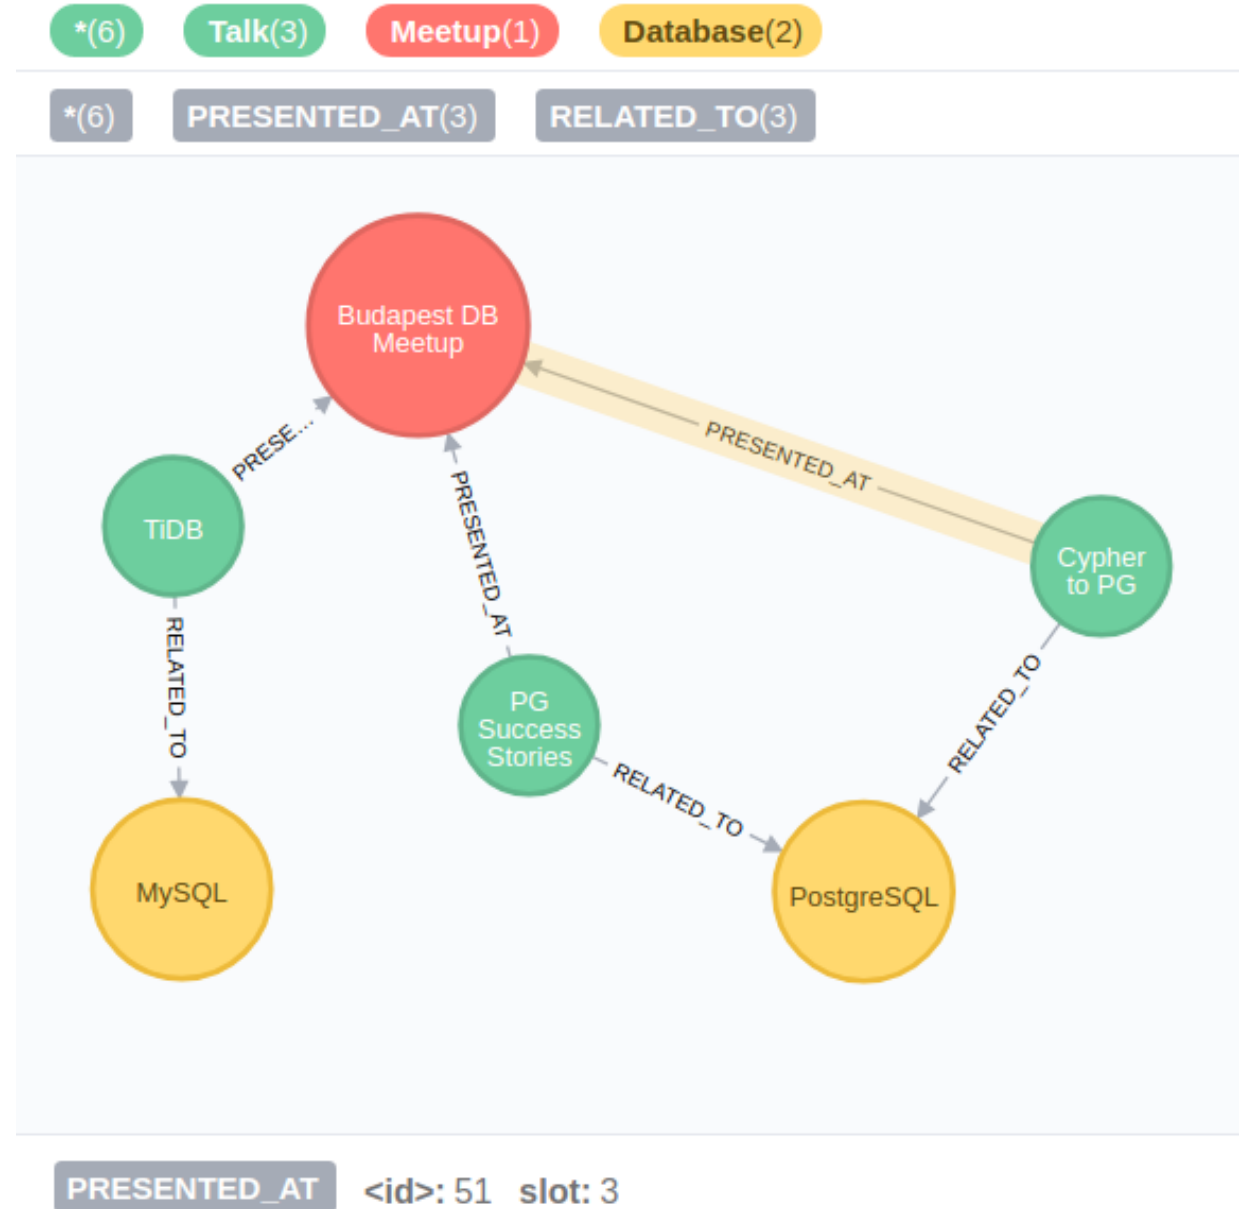
PROPERTY GRAPH DATABASES

NoSQL family


Data model:

- nodes
- edges
- properties

#1 query approach:
graph pattern matching



RANKINGS: POPULARITY CHANGES PER CATEGORY



**CYPHER
INTRODUCED
LATE 2013**

CYPHER AND OPENCYPHER

Cypher: query language of the Neo4j graph database.

„Cypher is a declarative, SQL-inspired language for describing patterns in graphs visually using an ascii-art syntax.”

MATCH

```
(d:Database)<-[:RELATED_TO]-(:Talk)-[:PRESENTED_AT]->(m:Meetup)
WHERE m.date = 'Tuesday, November 13, 2018'
RETURN d
```

„The openCypher project aims to deliver a full and open specification of the industry’s most widely adopted graph database query language: Cypher.” (late 2015)



openCypher

OPENCYIPHER SYSTEMS

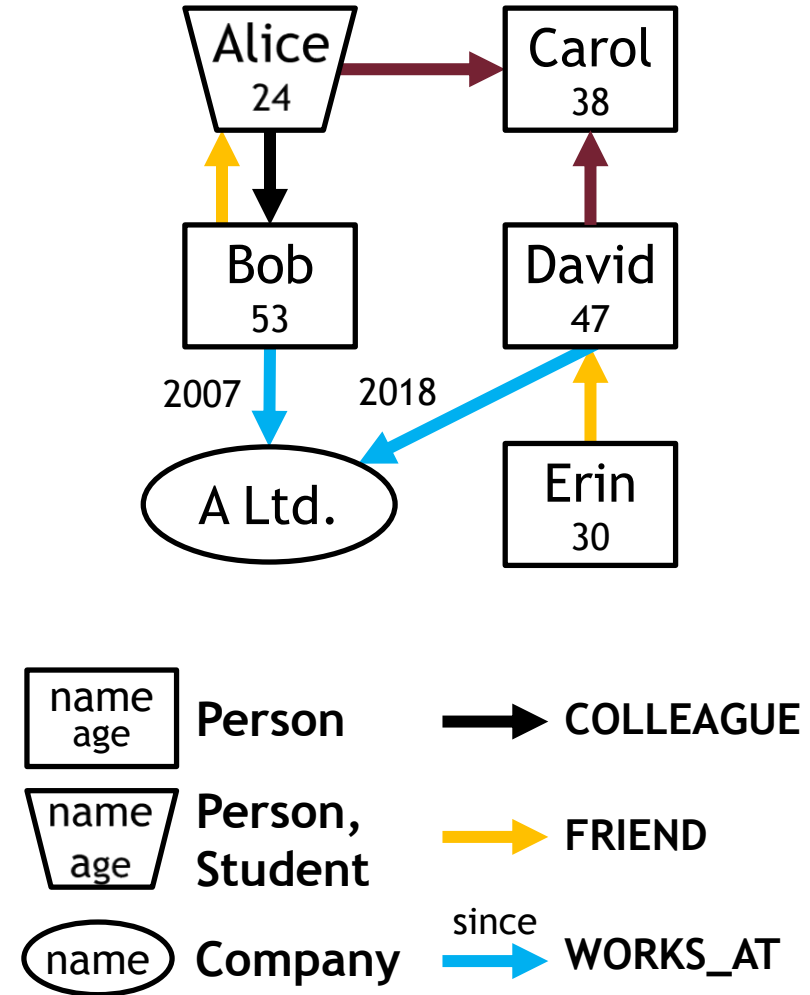
- Increasing adoption
- Relational databases:
 - SAP HANA
 - AGENS Graph
- Research prototypes:
 - Graphflow (Univesity of Waterloo, Canada)
 - ingraph (incremental graph engine @ BME)



(Source: Keynote talk @ GraphConnect NYC 2017)


PROPERTY GRAPHS

- Textbook graph: $G = (V, E)$
 - Dijkstra, Ford-Fulkerson, etc.
 - Homogenous nodes
 - Homogeneous edges
- Extensions
 - Labelled nodes
 - Typed edges
 - Properties
- The schema is implicit
- Very intuitive
 - *Things and connections*



GRAPH VS. RELATIONAL DATABASES

- Graph databases
 - Graph-based modelling is intuitive
 - Concise query language
- Relational databases
 - Most common
 - Many legacy systems
 - Efficient and mature
- No tools available to bridge the two
 - i.e. query data in RDBs as a graph
 - first you have to wrangle the graph out of the RDB

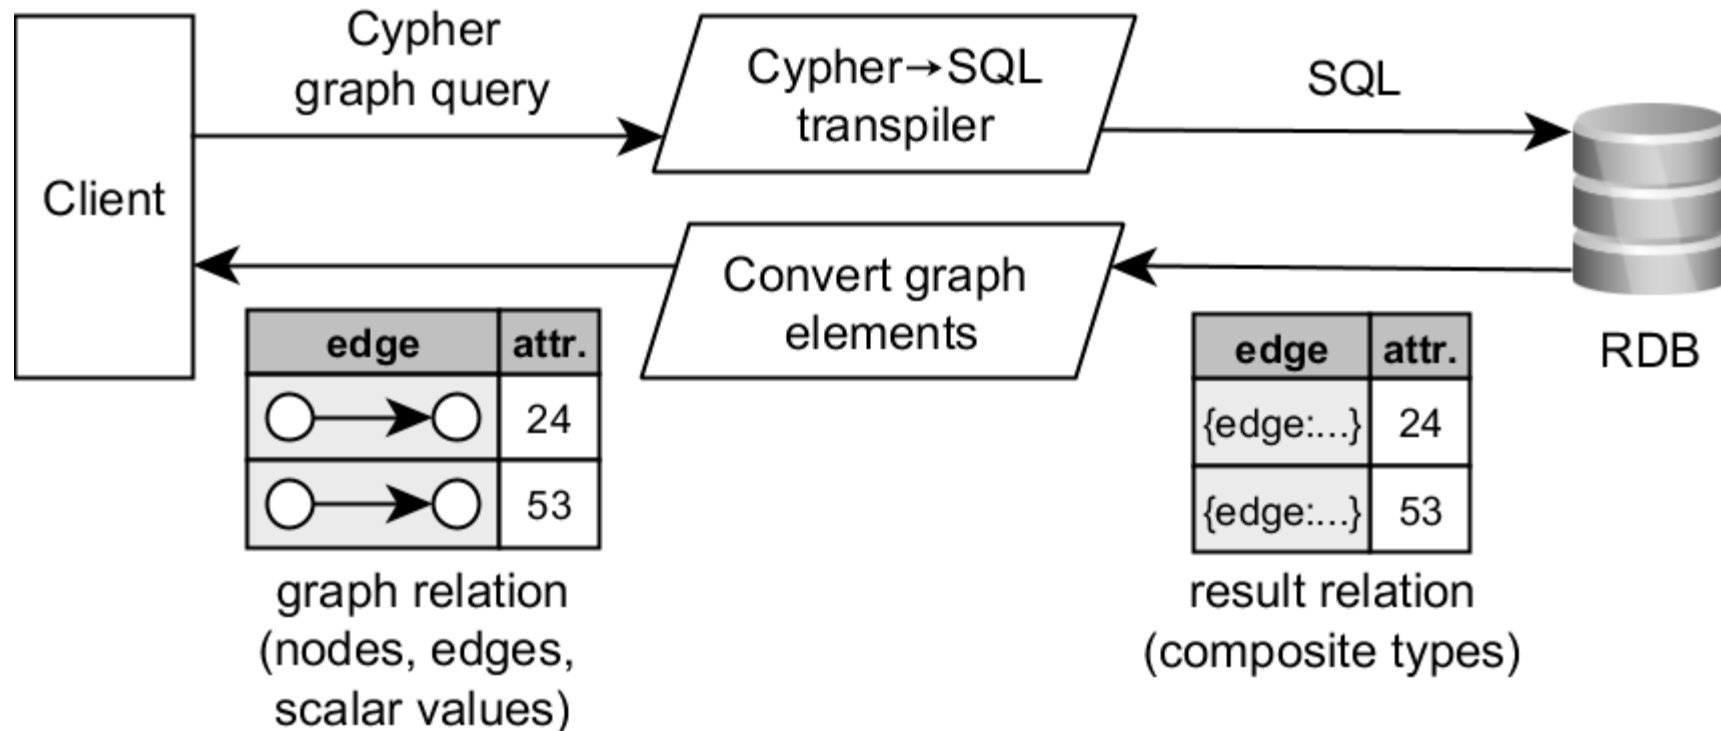


Tables	
Col1	Col2
1	A
2	B

PROPOSED APPROACH

To get the best of both worlds, map Cypher queries to SQL:

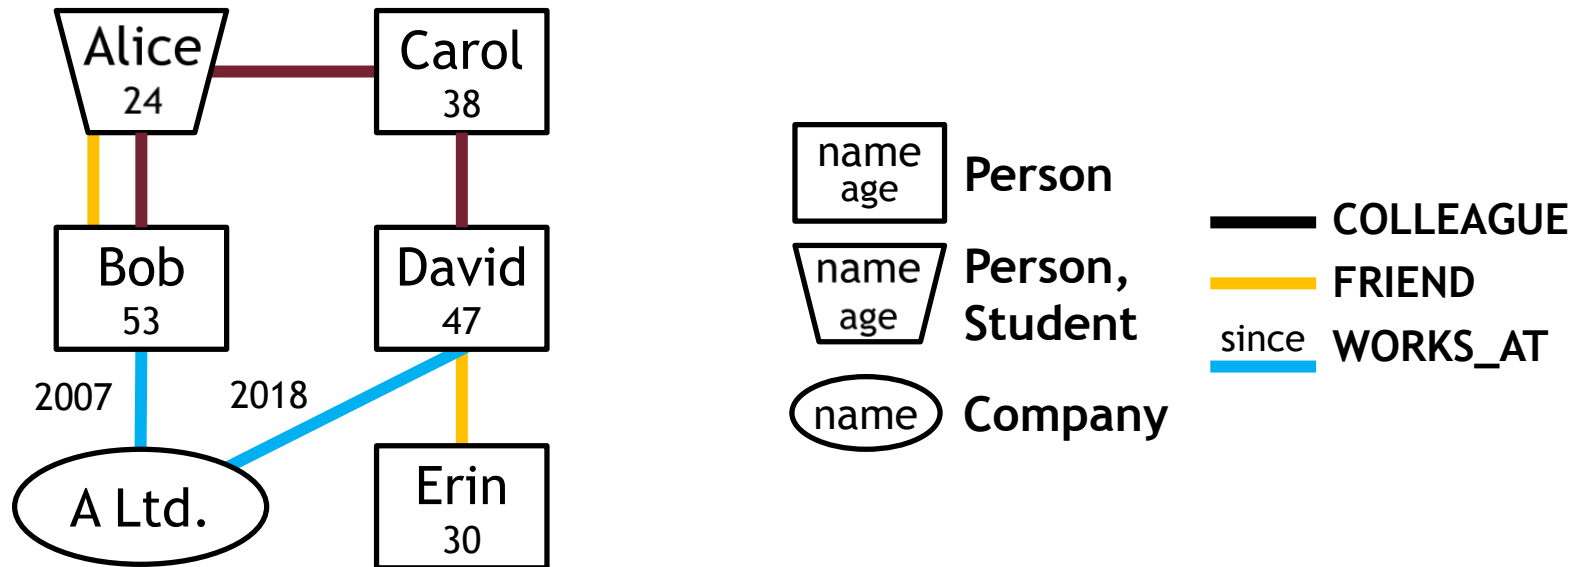
1. Formulate queries in Cypher
2. Execute inside an existing RDB
3. Return results as a graph relation



GRAPH QUERIES IN CYPHER

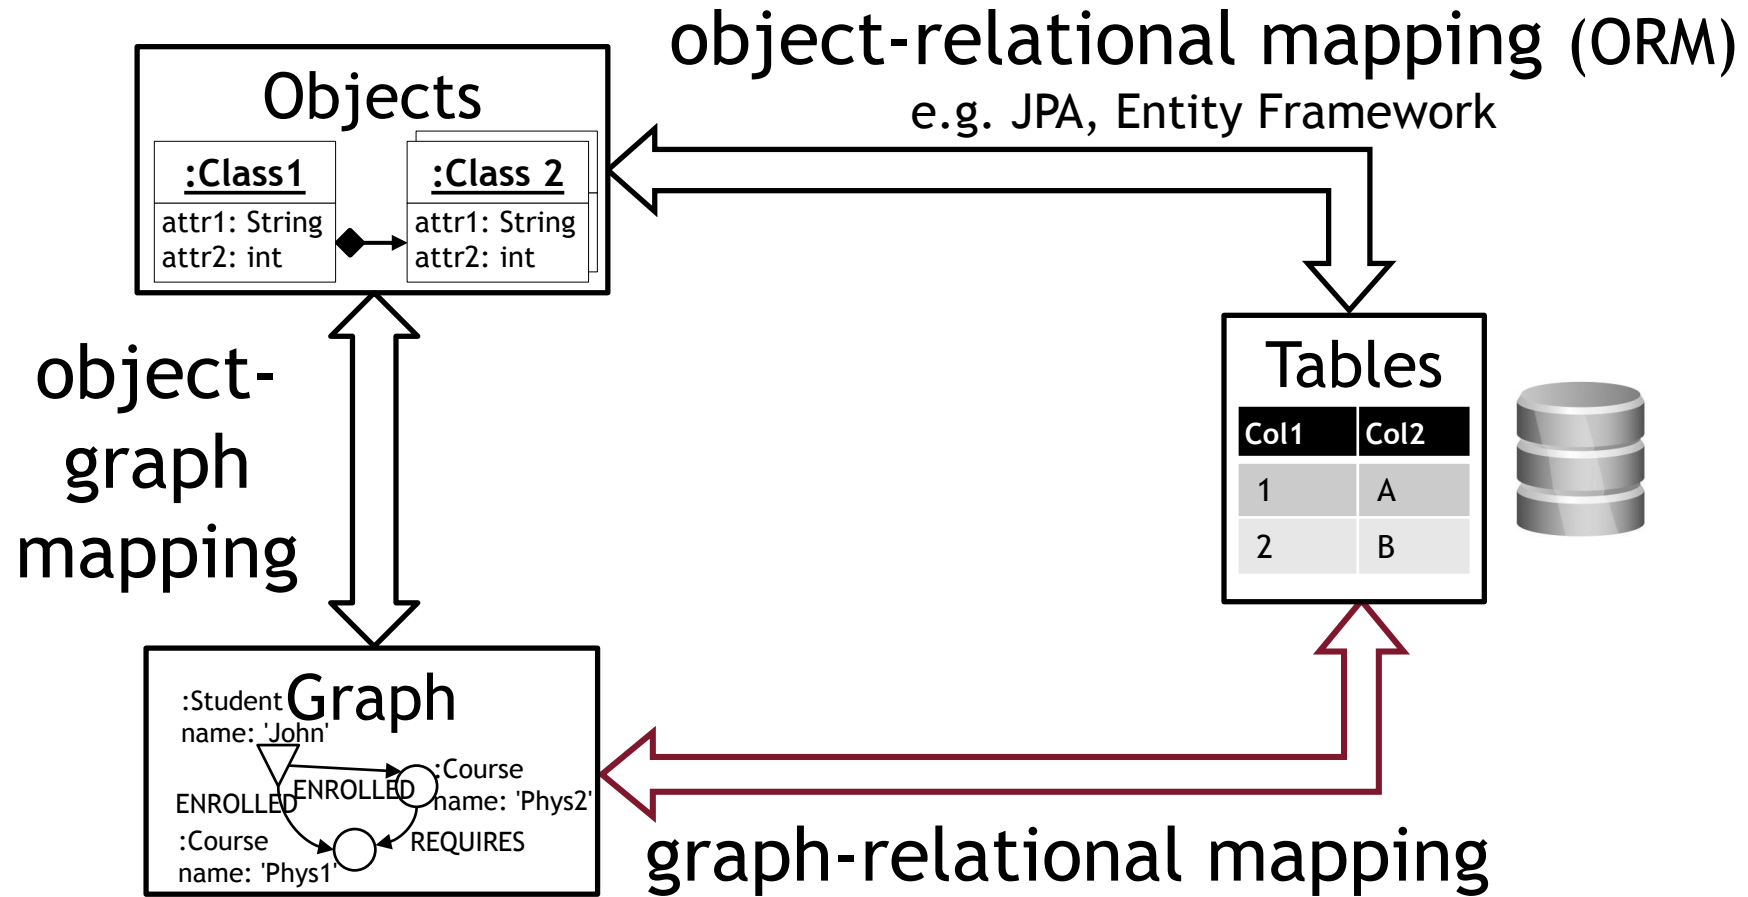
- Subgraph matching and graph traversals
- Example: Alice's colleagues and their colleagues

```
MATCH (p1:Person {name: 'Alice'})-[c:COLL*1..2]-(p2:Person)
RETURN p2.name
```



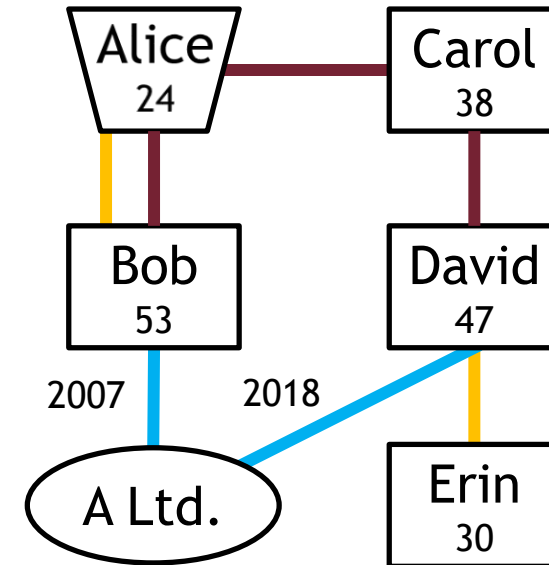
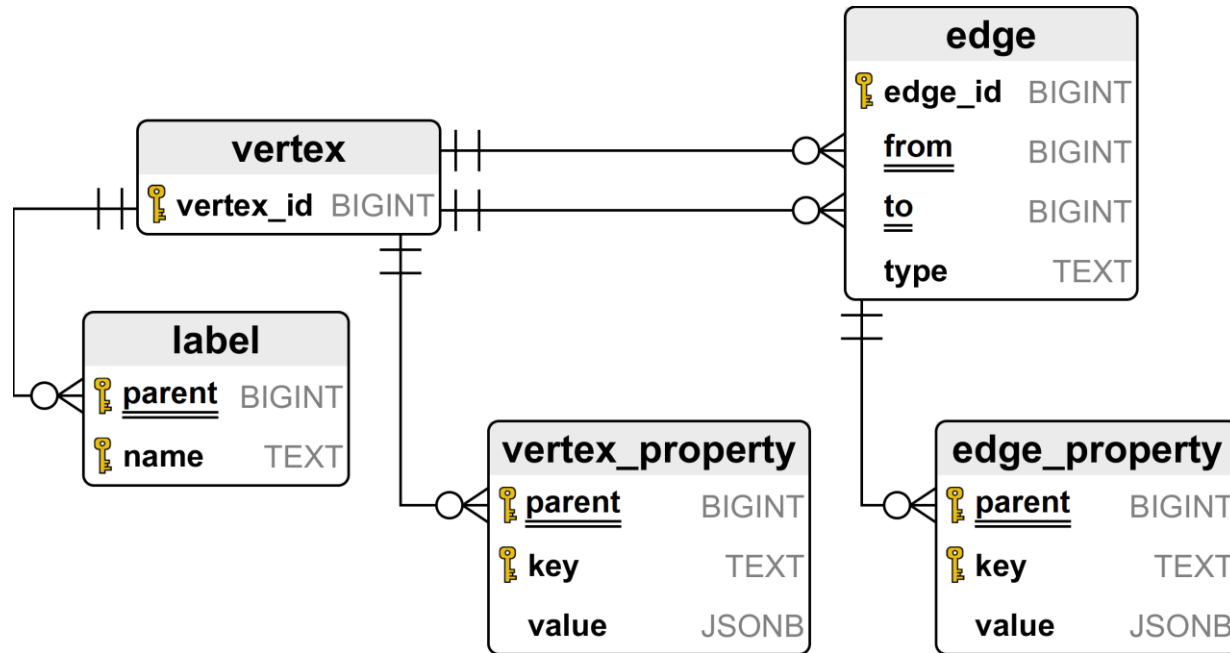
Data and Query Mapping

MAPPING BETWEEN DATA MODELS

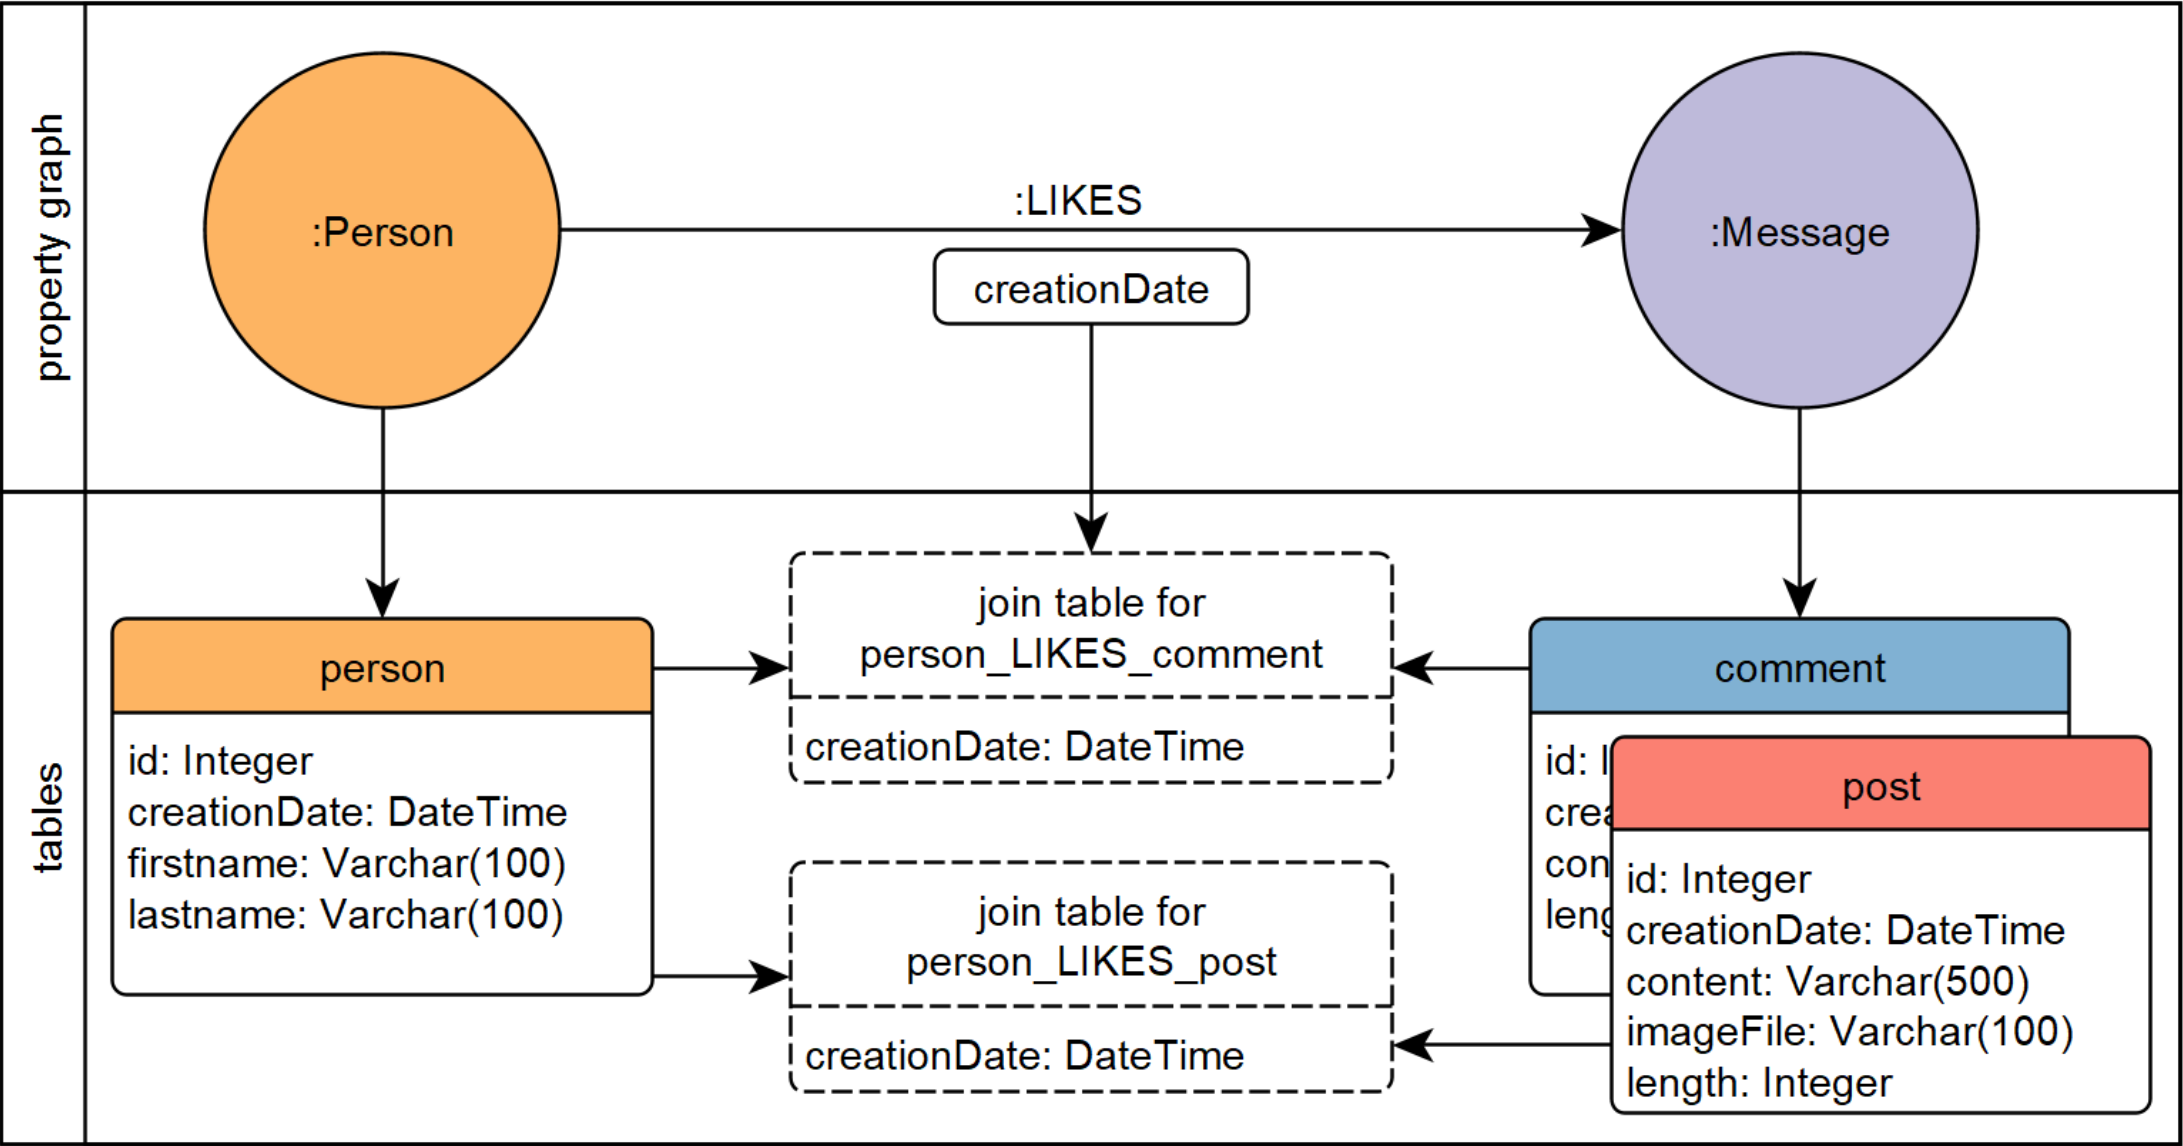


DATA MAPPING #1: GENERIC SCHEMA

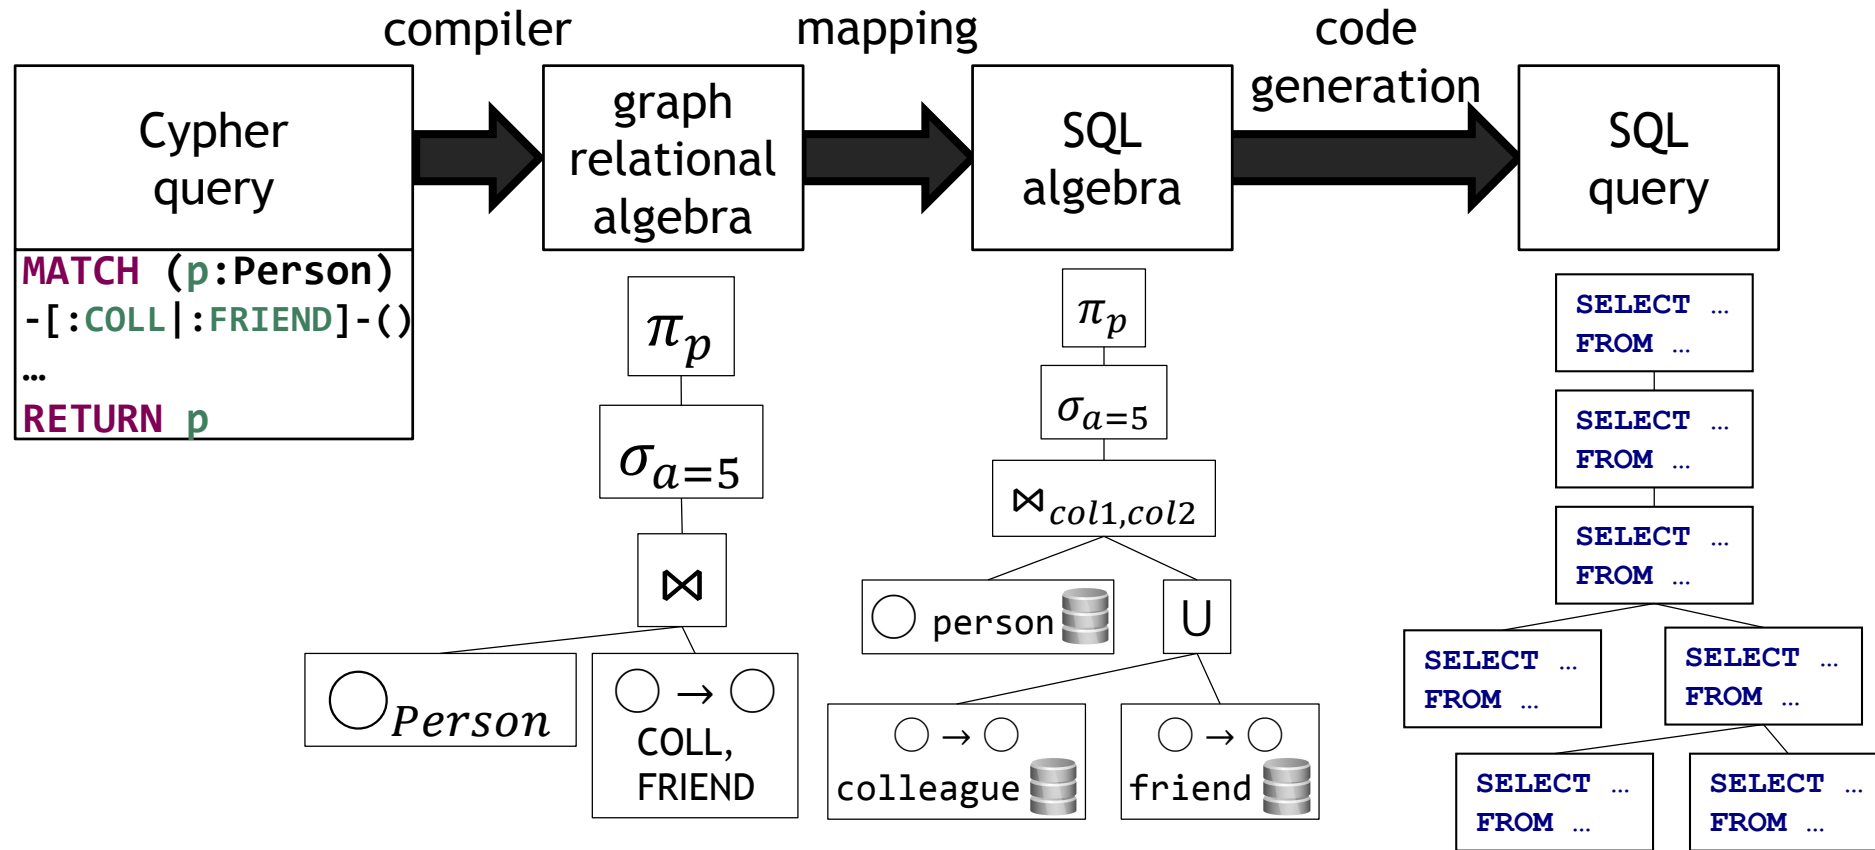
- Useful for representing schema-free data sets
- Hopelessly slow



DATA MAPPING #2: CONCRETE SCHEMA



QUERY MAPPING



Gábor Szárnyas, József Marton, Dániel Varró:
Formalising openCypher Graph Queries in Relational Algebra.
ADBIS 2017




A SIMPLE EXAMPLE

```
MATCH (node)
WITH node.name AS name
RETURN name
```

Cypher query

ingraph
mapping



```
SELECT "name"
FROM
(
  SELECT "node.name" AS "name"
  FROM
  (
    SELECT
      vertex_id AS "node",
      (SELECT value
       FROM vertex_property
       WHERE parent = vertex_id AND key = 'name') AS "node.name"
    FROM vertex
  )
)
```

CHALLENGES #1

- Variable length paths: union of multiple subqueries

```
MATCH (p1:Person {name: 'Alice'})-[c:COLL*1..2]-(p2:Person)
RETURN p2.name
```

- Unbound: WITH RECURSIVE (fixpoint-based evaluation)

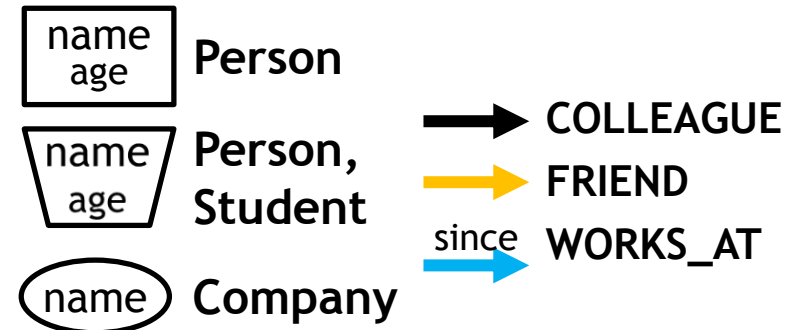

```
MATCH (p1:Person {name: 'Alice'})-[c:COLL*]-(p2:Person)
RETURN p2.name
```

- WITH RECURSIVE was introduced in SQL:1999 but
 - PostgreSQL 8.4+ (since 2009)
 - SQLite 3.8.3+ (since 2014)
 - MySQL 8.0.1+ (since 2017)

CHALLENGES #2

- Edges are directed
 - Undirectedness is modelled in the query
 - Union of both directions


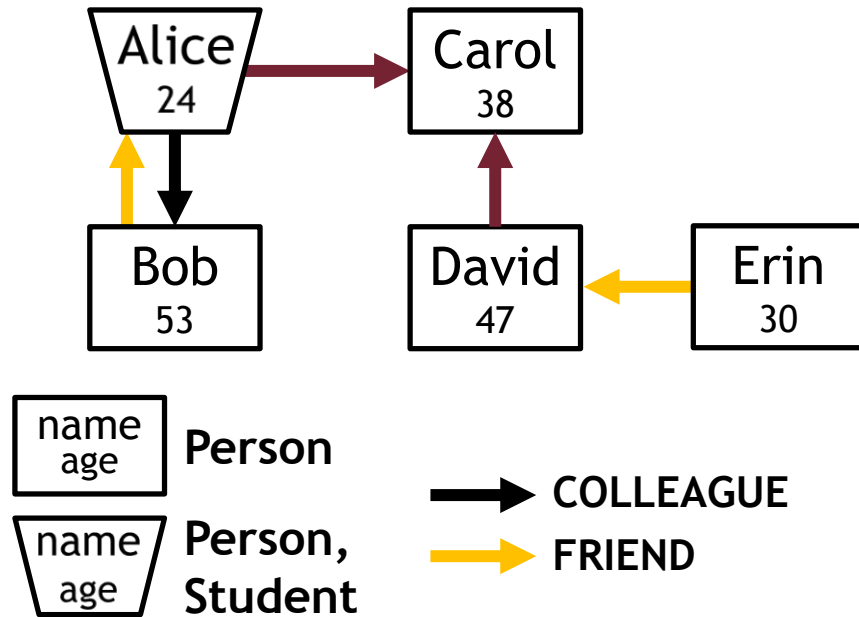
```
MATCH (p1:Person ...) -[:FRIEND]-(p2:Person),  
      (p2)-[:WORKS_AT]->(c:Company)  
RETURN p2.name, c.name
```



CHALLENGES #3

- Multiple tables as sources

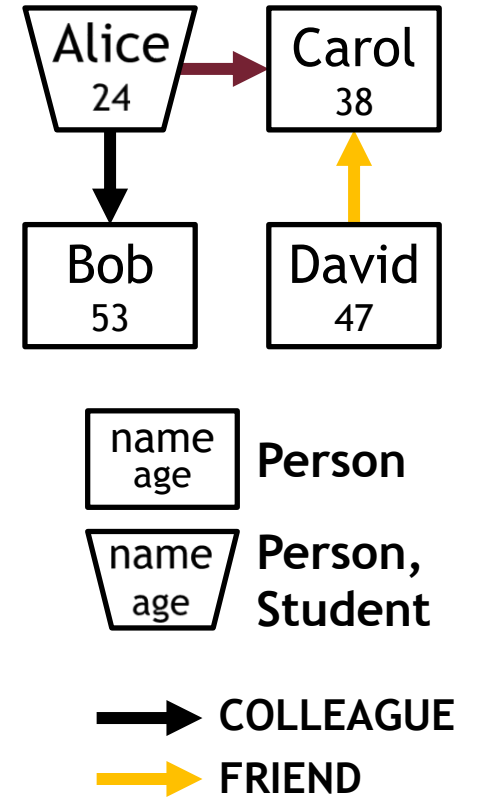
```
MATCH (p1:Person ...) -[:COLL|:FRIEND]-(p2:Person)
RETURN p2.name
```



CHALLENGES #1 #2 #3

Simple graph patterns turn to many subqueries

- Querying edges as undirected:
 - Enumerate edges in both directions ($2 \times$)
- Variable length paths ($1 \dots \mathbb{L}, *$):
 - Limited: enumerate $1, 2, \dots, L \rightarrow L \times$
 - Unlimited: use WITH RECURSIVE
- Multiple node labels/edge types:
 - Enumerate all source tables ($N \times$)



Total: union of $2 \times L \times N$ subqueries in SQL

```
MATCH (p1:Person ...) -[:COLL|:FRIEND*1..2]-(p2:Person)
```

A COMPLEX EXAMPLE

```
MATCH (:Person {id:$personId})-[:KNOWS]-(friend:Person)<-
[:HAS_CREATOR]-(message:Message)
WHERE message.creationDate <= $maxDate
RETURN
    friend.id AS personId,
    friend.firstName AS personFirstName,
    friend.lastName AS personLastName,
    message.id AS postOrCommentId,
```

```
CASE exists(message.content)
    WHEN true THEN message.content
    ELSE message.imageFile
END AS postOrCommentContent,
message.creationDate AS postOrCommentCreationDate
ORDER BY postOrCommentCreationDate DESC, toInteger(postOrCommentId) ASC
LIMIT 20
```

```
WITH
q0 AS
  (
    -- GetVerticesWithGTop
    SELECT
      ROW(0, p_personid)::vertex_type AS "_e186#0",
      "p_personid" AS "_e186.id#0"
    FROM person),
q1 AS
  (
    -- Selection
    SELECT * FROM q0 AS subquery
    WHERE ("_e186.id#0" = :personId)),
q2 AS
  (
    -- GetEdgesWithGTop
    SELECT ROW(0, edgeTable."k_personid")::vertex_type AS "_e186#0", ROW(0, edgeTable."k_person2id")::edge_type AS
    "_e187#0", ROW(0, edgeTable."k_person2id")::vertex_type AS "friend#2",
    toTable."p_personid" AS "friend.id#2", toTable."p_firstname" AS "friend.firstName#1", toTable."p_lastname" AS "friend.lastName#2"
    FROM "knows" edgeTable
    JOIN "person" toTable ON (edgeTable."k_person2id" = toTable."p_personid")),
q3 AS
  (
    -- GetEdgesWithGTop
    SELECT ROW(0, edgeTable."k_personid")::vertex_type AS "friend#2", ROW(0, edgeTable."k_personid", edgeTable."k_person2id")::edge_type AS
    "_e187#0", ROW(0, edgeTable."k_person2id")::vertex_type AS "_e186#0",
    fromTable."p_personid" AS "friend.id#2", fromTable."p_firstname" AS "friend.firstName#1", fromTable."p_lastname" AS "friend.lastName#2"
    FROM "knows" edgeTable
    JOIN "person" fromTable ON (fromTable."p_personid" = edgeTable."k_personid")),
q4 AS
  (
    -- UnionAll
    SELECT "_e186#0", "_e187#0", "friend#2", "friend.id#2", "friend.firstName#1", "friend.lastName#2" FROM q2
    UNION ALL
    SELECT "_e186#0", "_e187#0", "friend#2", "friend.id#2", "friend.firstName#1", "friend.lastName#2" FROM q3),
q5 AS
  (
    -- EquiJoinLike
    SELECT left_query."_e186#0", left_query."_e186.id#0", right_query."friend#2", right_query."friend.id#2", right_query."_e187#0",
    right_query."friend.lastName#2", right_query."friend.firstName#1" FROM
    q1 AS left_query
    INNER JOIN
    q4 AS right_query
    ON left_query."_e186#0" = right_query."_e186#0"),
q6 AS
  (
    -- GetEdgesWithGTop
    SELECT ROW(6, fromTable."m_messageid")::vertex_type AS "message#17", ROW(8, fromTable."m_messageid", fromTable."m_creatorid")::edge_type AS
    "_e188#0", ROW(0, fromTable."m_creatorid")::vertex_type AS "friend#2",
    fromTable."m_messageid" AS "message.id#2", fromTable."m_content" AS "message.content#2", fromTable."m_ps_imagefile" AS
    "message.imageFile#0", fromTable."m_creationdate" AS "message.creationDate#13"
    FROM "message" fromTable
    WHERE (fromTable."m_c_replyof" IS NULL)),
```


```
q7 AS
  (
    -- GetEdgesWithGTop
    SELECT ROW(6, fromTable."m_messageid")::vertex_type AS "message#17", ROW(8, fromTable."m_messageid", fromTable."m_creatorid")::edge_type AS
    "_e188#0", ROW(0, fromTable."m_creatorid")::vertex_type AS "friend#2",
    fromTable."m_messageid" AS "message.id#2", fromTable."m_content" AS "message.content#2", fromTable."m_creationdate" AS
    "message.creationDate#13"
    FROM "message" fromTable
    WHERE (fromTable."m_c_replyof" IS NOT NULL)),
q8 AS
  (
    -- UnionAll
    SELECT "message#17", "_e188#0", "friend#2", "message.id#2", "message.content#2", "message.imageFile#0", "message.creationDate#13" FROM q6
    UNION ALL
    SELECT "message#17", "_e188#0", "friend#2", "message.id#2", "message.content#2", NULL AS "message.imageFile#0", "message.creationDate#13"
    FROM q7),
q9 AS
  (
    -- EquiJoinLike
    SELECT left_query."_e186#0", left_query."_e186.id#0", left_query."_e187#0", left_query."friend#2", left_query."friend.id#2",
    left_query."friend.firstName#1", left_query."friend.lastName#2", right_query."message#17", right_query."message.id#2",
    right_query."message.imageFile#0", right_query."_e188#0", right_query."message.creationDate#13", right_query."message.content#2" FROM
    q5 AS left_query
    INNER JOIN
    q8 AS right_query
    ON left_query."friend#2" = right_query."friend#2"),
q10 AS
  (
    -- AllDifferent
    SELECT * FROM q9 AS subquery
    WHERE is_unique(ARRAY[::edge_type[] || "_e188#0" || "_e187#0"])),
q11 AS
  (
    -- Selection
    SELECT * FROM q10 AS subquery
    WHERE ("message.creationDate#13" <= :maxDate)),
q12 AS
  (
    -- Projection
    SELECT "friend.id#2" AS "personId#0", "friend.firstName#1" AS "personFirstName#0", "friend.lastName#2" AS "personLastName#0",
    "message.id#2" AS "postOrCommentId#0", CASE WHEN ("message.content#2" IS NOT NULL = true) THEN "message.content#2"
    ELSE "message.imageFile#0"
    END AS "postOrCommentContent#0", "message.creationDate#13" AS "postOrCommentCreationDate#0"
    FROM q11 AS subquery),
q13 AS
  (
    -- SortAndTop
    SELECT * FROM q12 AS subquery
    ORDER BY "postOrCommentCreationDate#0" DESC NULLS LAST, ("postOrCommentId#0")::BIGINT ASC NULLS FIRST
    LIMIT 20)
SELECT "personId#0" AS "personId", "personFirstName#0" AS "personFirstName", "personLastName#0" AS "personLastName", "postOrCommentId#0" AS
"postOrCommentId", "postOrCommentContent#0" AS "postOrCommentContent", "postOrCommentCreationDate#0" AS "postOrCommentCreationDate"
FROM q13 AS subquery
```

Benchmarks

BENCHMARKS: LINKED DATA BENCHMARK COUNCIL

LDBC is a non-profit organization dedicated to establishing benchmarks, benchmark practices and benchmark results for graph data management software.

The Social Network Benchmark is an industrial and academic initiative, formed by principal actors in the field of graph-like data management.




LDBC IN A NUTSHELL



Peter Boncz, Thomas Neumann, Orri Erling,
TPC-H Analyzed: Hidden Messages and Lessons Learned from an Influential Benchmark,
TPCTC 2013



Orri Erling et al.,
The LDBC Social Network Benchmark: Interactive Workload,
SIGMOD 2015

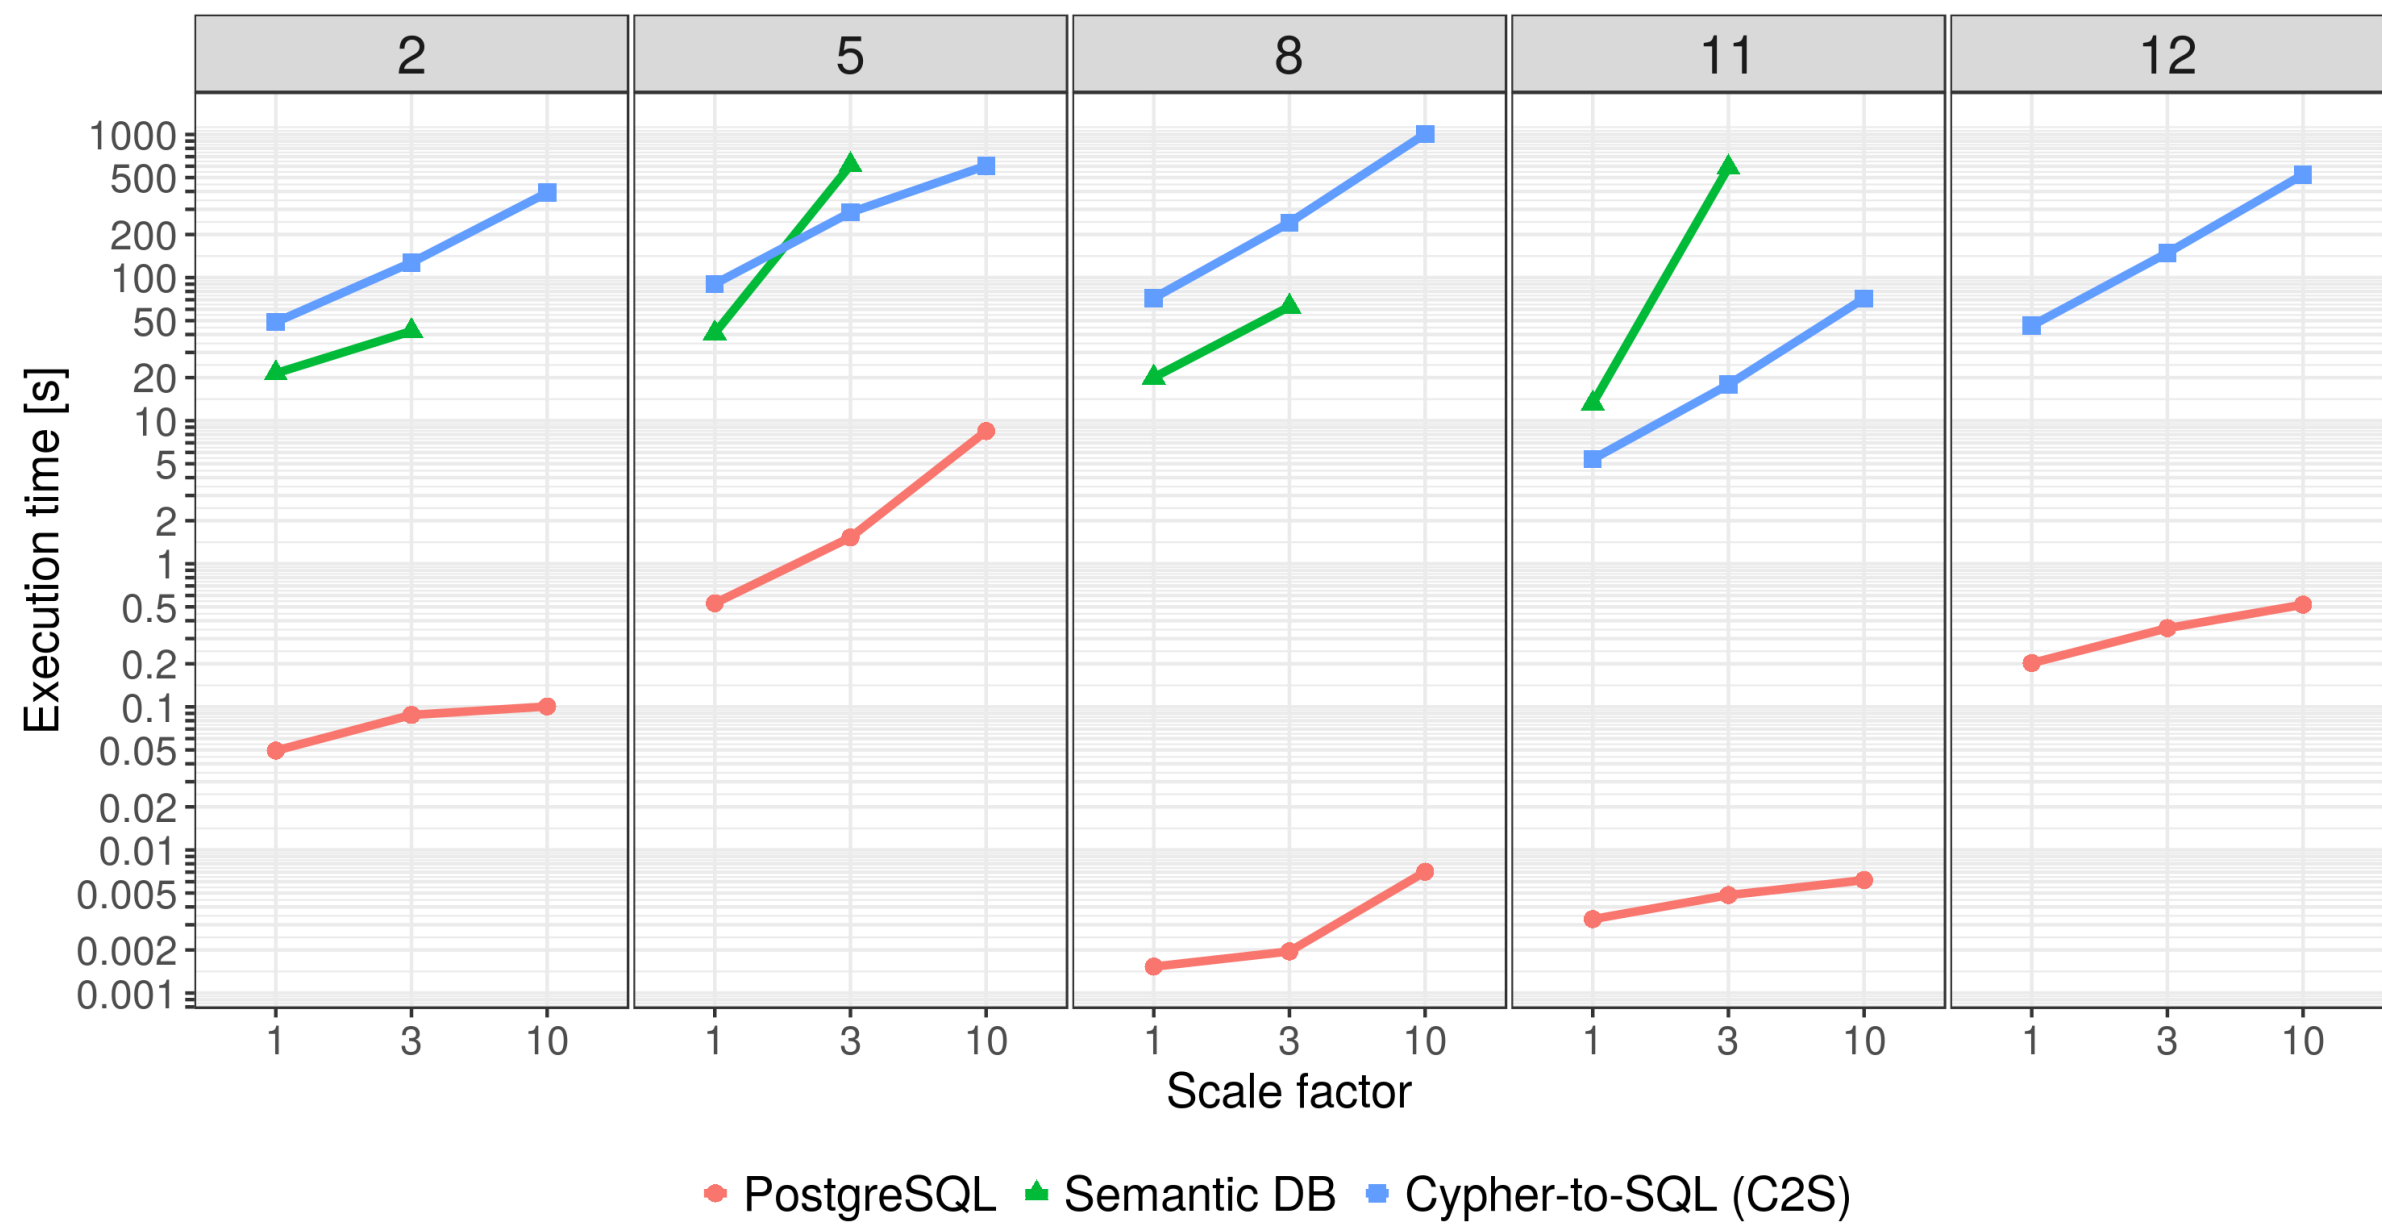


Gábor Szárnyas, József Marton, János Benjamin Antal et al.:
An early look at the LDBC Social Network Benchmark's BI Workload.
GRADES-NDA at SIGMOD, 2018

PERFORMANCE EXPERIMENTS

- LDBC Interactive workload
- Tools
 - PostgreSQL (reference implementation)
 - Cypher-to-SQL queries on PostgreSQL
 - Semantic database (anonymized)
- Geometric mean of 20+ executions

BENCHMARK RESULTS ON LDBC QUERIES



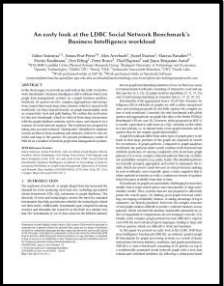
RELATED PROJECTS

- Cypher for Apache Spark
 - Neo4j's project
 - Executes queries in Spark
 - Read-only
- Cytosm
 - Cypher to SQL Mapping
 - HP Labs for Vertica
 - Project abandoned in 2017
 - gTop (graph topology) reused

Tool	Source	Target	OSS	Updates	Paths
CAPS	Cypher	SparkSQL	✓	✗	✗
Cytosm	Cypher	Vertica SQL	✓	✗	✗
Cypher-to-SQL	Cypher	PostgreSQL	✓	✓	✓

SUMMARY

- Mapping property graph queries to SQL is challenging
 - Similar to ORM
 - + edge properties
 - + reachability
- Initial implementation: C2S
 - Moderate feature coverage
 - Poor performance
 - Needs some tweaks, e.g. work around CTE optimization fences



Gábor Szárnyas, József Marton, János Maginecz, Dániel Varró:
Incremental View Maintenance on Property Graphs.
arXiv preprint 2018

RELATED RESOURCES

ingraph and C2S

github.com/ftsrg/ingraph

Cypher for Apache Spark

github.com/opencypher/cypher-for-apache-spark

Cytosm

github.com/cytosm/cytosm

LDBC

github.com/ldbc/

Thanks for the contributions to the whole ingraph team.

